

Structuring Machine Learning Projects

Week 1:

1. Chain of assumption in ML:
 - a. Fit training set well on cost function.(Human level performance)
 - i. Bigger network or optimization
 - b. Fit dev set well on cost function
 - i. Regularize or bigger training set
 - c. Fit test set well on cost function
 - i. Bigger dev set
 - d. Performs well in real world
 - i. Change dev set or cost function
2. Single number eval metric:
 - a. So if classifier A has 95% precision, this means that when classifier A says something is a cat, there's a 95% chance it really is a cat. And recall is, of all the images that really are cats, what percentage were correctly recognized by your classifier? So what percentage of actual cats, Are correctly recognized?
 - b. So if classifier A is 90% recall, this means that of all of the images in, say, your dev sets that really are cats, classifier A accurately pulled out 90% of them.
 - c. F1 score – Average of precision and recall(Harmonic mean)
3. Train/dev/test distribution:
 - a. Dev and test must come from the same distribution
 - b. Choose a dev set and test set to reflect your data you expect to get in future and consider important to well on
4. Size of the dev and test sets:
 - a. Set your test set to be big enough to give high confidence in the overall performance of your system
5. When to change dev/test set and metrics:
 - a. But the high level of take away is, if you find that evaluation metric is not giving the correct rank order preference for what is actually better algorithm,
6. Orthogonalization:
 - a. Define metrics
 - b. Worry separately about how to do well on this metric
7. The two fundamental assumption of supervised learning:
 - a. You can fit the training set pretty well
 - b. The training set performance generalizes pretty well to the dev/test set
8. Reducing avoidable bias and variance:
 - a. Reduce avoidable bias:
 - i. Train big model
 - ii. Train longer or better optimization algorithms
 - iii. Find better set of hyperparameters
 - b. Reduce variance:
 - i. More data
 - ii. Regularization (L2, dropout, data augmentation)
 - iii. Search better hyperparameters

Week 2:

9. Correcting incorrect dev/test set examples:

- a. Apply some process to your dev and test sets to make sure they come from the same distribution
 - b. Consider examining examples of your algorithm got right as well as ones it got wrong
- 10. Build your system quickly, then iterate:
 - a. Set up dev/test set and metric
 - b. Build initial system quickly
 - c. Use Bias/Variance analysis and error analysis to prioritize next steps
- 11. Addressing data mismatch:
 - a. Carry out manual error analysis to try to understand difference between training and dev/test sets
 - b. Make training data more similar or collect more data similar to dev/test sets
- 12. Transfer learning:
 - a. Transfer learning is a machine learning technique where a model trained on one task is re-purposed on a second related task.
 - b. Transfer learning makes sense when you have a lot of data for the problem you're transferring from and usually relatively less data for the problem you're transferring to.
- 13. Multi-task learning:
 - a. Training on a set of tasks that could benefit from having shared low-level features
 - b. Amount of data you have for each task is quite similar
 - c. Enables you to train an NN doing with many tasks
- 14. Pros and cons of End-to-end deep learning:
 - a. Pros:
 - i. Let data speak
 - ii. Less hand-designing of components needed
 - b. Cons:
 - i. Need a lot of data
 - ii. Exclude potentially useful hand-designing things
 - c. Applying E2E:
 - i. Key question: Do you have sufficient data to learn a function of the complexity needed to map x to y